

SINDHU S

GenAI Engineer

✉ sindhufeb99@gmail.com
☎ +91-7339604736
📍 Chennai, India

TECHNICAL SKILLS

GenAI & LLMs

LangChain, OpenAI API,
RAG Pipelines, FAISS,
Pinecone, Prompt
Engineering, Hugging Face,
Streamlit, AWS
Bedrock, Open Search, Open
AI

Data Engineering

Snowflake, ELT Pipelines,
Data Modelling, Snowflake
Streams & Tasks

Cloud

AWS (S3), Snowflake Cloud

Programming

Python (pandas, NumPy,
FastAPI), SQL, JavaScript

Databases

Snowflake, MySQL, FAISS,
Pinecone (Vector DB)

Tools

GitHub Desktop, Docker,
Jupyter Notebook

EDUCATION

MSc (Computer Science)

Government Arts College
Salem-7

PROFESSIONAL SUMMARY

GenAI Engineer with 4+ years hands-on experience in building scalable AI-driven and data solutions, backed by relevant experience in Data Engineering. Started as a Data Engineer, developing Snowflake-based data pipelines, ELT workflows, and data warehouse systems for structured financial data. Transitioned into Generative AI, with hands-on experience in building RAG-based applications, LLM integrations, and AI-powered solutions for document analysis and insights. Strong understanding of the end-to-end data lifecycle—from data ingestion and transformation to delivering business-focused AI insights.

PROFESSIONAL EXPERIENCE

Company: ENTRUST SOFTWARE & IT SERVICES

06/2023 – Present | Chennai

GenAI Data Engineer

Client: NovaTrust

- Architected and delivered a **production-grade GenAI platform** for analyzing financial reports and pitch books, reducing manual analysis effort and accelerating decision-making
- Engineered scalable **RAG pipelines with vector search**, significantly improving accuracy and contextual relevance of information retrieval from unstructured data
- Built robust **document ingestion and processing pipelines** (parsing, chunking, embeddings, indexing) to handle large-scale financial documents efficiently
- Designed and optimized **LLM prompt strategies** for high-quality summarization, risk detection, and financial insight generation
- Developed **high-performance REST APIs** enabling natural language querying, driving measurable gains in analyst productivity and turnaround time
- Ensured **scalability and reliability** through modular architecture and optimized data workflows

02/2022 – 05/2023 | Chennai

- Developed a **RAG-based system** that converts user queries into SQL, retrieves relevant trade data from Snowflake, and generates human-readable responses using LLMs
- Designed and implemented data pipelines to ingest, cleanse, and structure high-volume financial transaction data for AI consumption
- Created embeddings for financial datasets and metadata to enable semantic search across trades, portfolios, and historical transactions
- Integrated LLMs (via APIs like Groq/OpenAI) with LangChain to orchestrate query understanding, context retrieval, and response generation
- Enabled near real-time insights by combining Snowflake Streams/Tasks with AI-driven query execution workflows
- Optimized query performance and context retrieval to ensure low-latency responses for interactive chatbot use cases